

# System Architecture Overview

The project is a distributed file storage system implemented in C, consisting of four servers (s1.c, s2.c, s3.c, s4.c) and a client (w25clients.c):

- S1 (Port 9086): The primary server that handles .c files locally and forwards .pdf, .txt, and .zip files to S2, S3, and S4, respectively. It coordinates client requests and communicates with secondary servers.
- S2 (Port 9087): Manages .pdf files, supporting upload, download, removal, listing, and tar archiving.
- S3 (Port 9088): Manages .txt files, with similar functionality to S2.
- S4 (Port 9089): Manages .zip files, mirroring S2 and S3's operations.
- Client (w25clients.c): Connects to S1 to perform file operations: upload (uploadf), download (downlf), remove (removef), list files (dispfnames), and download tar archives (downltar).

Each server uses TCP sockets for reliable communication, stores files in designated directories (./S1, ./S2, ./S3, ./S4), and maintains a log file (S1.log, S2.log, etc.). The system is designed to mimic a distributed storage system, with S1 acting as a gateway and secondary servers specializing in specific file types.

## Libraries Used Across All Files

- Standard Libraries:
  - <stdio.h>: For I/O operations (printf, fopen, fprintf, fclose).
  - <stdlib.h>: For memory management and exit (malloc, exit).
  - <string.h>: For string manipulation (strcmp, strncpy, snprintf, strstr).
- System and Networking:
  - <unistd.h>: For system calls (close, read, write, fork, usleep, access).
  - <sys/socket.h>: For socket operations (socket, bind, listen, accept, connect, send, setsockopt).
  - <netinet/in.h>: For Internet address structures (struct sockaddr\_in).
  - <arpa/inet.h>: For IP address conversion (inet\_pton, inet\_ntop, htons).
  - <sys/stat.h>: For file and directory operations (stat, mkdir, fstat).
  - <fcntl.h>: For file control (open, O\_WRONLY, O\_CREAT, O\_RDONLY).
  - <sys/wait.h>: For process management in S1 (waitpid).
  - <sys/time.h>: For socket timeouts (struct timeval).
- Directory and Time:
  - <dirent.h>: For directory operations (opendir, readdir, closedir).
  - <time.h>: For timestamps (time, ctime).
  - <errno.h>: For error details (strerror, errno).

## Constants

- Ports: PORT\_S1 (9086), PORT\_S2 (9087), PORT\_S3 (9088), PORT\_S4 (9089).
  - Buffer Size: BUFFER\_SIZE (4096 bytes) for data transfers.
  - Path Length: MAX\_PATH (1024) for file paths.
  - Directories: S1\_DIR (./S1), S2\_DIR (./S2), S3\_DIR (./S3), S4\_DIR (./S4).
  - Log Files: S1.log, S2.log, S3.log, S4.log.
  - Server IPs (in s1.c): S2\_IP, S3\_IP, S4\_IP as 127.0.0.1 (localhost).
  - Client: SERVER\_IP (127.0.0.1), PORT\_S1 (9086).
-

## Function Explanations

Functions Common to s1.c, s2.c, s3.c, s4.c

These functions are nearly identical across servers, differing only in directory names (S1\_DIR, S2\_DIR, etc.) and file types in list\_files and tar\_files.

### 1. log\_message(const char \*msg)

- Purpose: Logs a message with a timestamp to the server's log file (S1.log, S2.log, etc.).
- Implementation:
  - Opens the log file in append mode using fopen(LOG\_FILE, "a").
  - If fopen fails (e.g., disk full, permissions), silently returns to keep the server running, prioritizing robustness over log loss.
  - Gets the current time with time(&now) and converts to a string using ctime.
  - Removes the trailing newline from the time string by setting time\_str[strlen(time\_str) - 1] = '\0'.
  - Writes the message in the format [timestamp] msg\n using fprintf.
  - Closes the file with fclose.
  - Example log entry: [Mon Apr 14 10:00:00 2025] Server started.
- Libraries Used:
  - <stdio.h>: fopen, fprintf, fclose.
  - <time.h>: time, ctime.
  - <string.h>: strlen.
- Error Handling:
  - Silent failure on fopen to avoid crashing.
- Edge Cases:
  - Disk full: Log entry is skipped.
  - Long messages: Handled by fprintf, no truncation risk within typical log message lengths.
- Performance Note:
  - Opening/closing the file per log is I/O-intensive. Buffering or keeping the file open could optimize high-frequency logging (not implemented).
- Viva Tips:
  - Explain logging's role in debugging distributed systems, tracing S1's interactions with S2–S4.
  - Mention the silent failure trade-off: ensures server uptime but risks missing logs.
- Example:
  - Input: log\_message("New connection").
  - Output in S1.log: [Mon Apr 14 10:00:01 2025] New connection.

### 2. receive\_file(int socket, const char \*filepath)

- Purpose: Receives a file from a socket and saves it to filepath.
- Implementation:
  - Sets a 5-second receive timeout using setsockopt(socket, SOL\_SOCKET, SO\_RCVTIMEO, &tv, sizeof(tv)) with tv.tv\_sec = 5.
  - Sends "READY" (5 bytes) using send(socket, "READY", 5, 0) to signal readiness.
  - Logs "Sent READY signal to client".
  - Reads the file size (off\_t file\_size) using read(socket, &file\_size, sizeof(off\_t)).
    - If read fails (e.g., connection closed, timeout), logs the error with snprintf and strerror, sends "ERROR: Size not received", and returns.

- Opens the file using `open(filepath, O_WRONLY | O_CREAT | O_TRUNC, 0666)`.
    - If open fails (e.g., no permissions), logs the error, sends "ERROR: File creation failed", and returns.
  - Reads data in chunks (`BUFFER_SIZE = 4096`) using read and writes to the file using write.
    - Tracks bytes received (received).
    - If read fails (e.g., client disconnects), logs and sends "ERROR: Read failed".
    - If write fails (e.g., disk full), logs and sends "ERROR: Write failed".
  - Closes the file with `close(fd)`.
  - If received == `file_size`, sends "SUCCESS: File uploaded"; else, sends "ERROR: Incomplete transfer" and logs details.
  - Resets timeout to 0 using `setsockopt`.
  - Logs key steps: READY signal, file size received, success/failure.
- Libraries Used:
  - `<sys/socket.h>`: send, read, `setsockopt`.
  - `<fcntl.h>`: open, write, close.
  - `<sys/time.h>`: struct `timeval`.
  - `<string.h>`: `snprintf`.
  - `<errno.h>`: `strerror`.
- Error Handling:
  - Handles socket errors (timeout, disconnect), file open/write failures, and incomplete transfers.
  - Sends descriptive error messages to the client.
- Edge Cases:
  - Empty File: If `file_size == 0`, creates an empty file and sends "SUCCESS".
  - Large Files: Loops handle any size, as chunks are 4096 bytes.
  - Invalid Path: Assumes caller validates `filepath` (starts with `S1_DIR`).
- Security:
  - Relies on caller to ensure `filepath` is within server directory, preventing writes outside (e.g., `/etc/passwd`).
- Performance:
  - Buffer size (4096) balances memory and I/O efficiency, common for network transfers.
- Viva Tips:
  - Highlight the READY-size-data protocol as a handshake for reliable transfers.
  - Discuss why timeout is critical (prevents server hanging on slow clients).
  - Explain buffer size choice (standard for TCP transfers).
- Example:
  - Input: `receive_file(sock, "./S1/dir/file.c")`.
  - Process: Receives size, saves data, sends "SUCCESS: File uploaded".

### **3. `send_file(int socket, const char *filepath)`**

- Purpose: Sends a file over a socket to the client.
- Implementation:
  - Opens the file using `open(filepath, O_RDONLY)`.
    - If open fails (e.g., file not found), logs "File open failed", sends `file_size = -1`, sends "ERROR: File not found", and returns.
  - Gets file size using `fstat(fd, &st)` and `st.st_size`.
  - Sends "READY" using `send(socket, "READY", 5, 0)`.

- Waits 10ms with usleep(10000) to ensure client readiness.
  - Sends file size (off\_t) using send(socket, &file\_size, sizeof(off\_t), 0).
  - Reads file in chunks (BUFFER\_SIZE) using read and sends using send.
    - Continues until read returns 0 (end of file).
  - Closes file with close(fd).
  - Logs "File sent successfully" on completion.
- Libraries Used:
  - <fcntl.h>: open, read, close.
  - <sys/stat.h>: fstat, struct stat.
  - <sys/socket.h>: send.
  - <unistd.h>: usleep.
- Error Handling:
  - Handles file open failure with clear client feedback.
  - Assumes send succeeds (TCP retries internally).
- Edge Cases:
  - Empty File: Sends file\_size = 0, handled correctly by client.
  - File Deleted: Sends file\_size = -1 and error message.
- Performance:
  - usleep delay is heuristic; could be tuned for network conditions.
- Viva Tips:
  - Explain TCP's reliability vs. UDP (no packet loss handling needed).
  - Discuss usleep as a sync mechanism, potential for removal with better client coordination.
- Example:
  - Input: send\_file(sock, "./S1/dir/file.c").
  - Output: Sends size and data, logs success.

## 4. **create\_directory(const char \*path)**

- Purpose: Creates a directory and its parents recursively.
- Implementation:
  - Copies path to a temporary buffer using snprintf(tmp, MAX\_PATH, "%s", path).
  - Starts parsing after base directory (S1\_DIR, S2\_DIR, etc.) using p = tmp + strlen(S1\_DIR) + 1.
  - Iterates through tmp, finding slashes (/).
    - For each slash, null-terminates (\*p = '\0'), calls mkdir(tmp, 0777), restores slash (\*p = '/').
  - Creates final directory with mkdir(tmp, 0777).
  - Logs "Directory created or exists".
- Libraries Used:
  - <string.h>: snprintf, strlen.
  - <sys/stat.h>: mkdir.
- Error Handling:
  - mkdir fails silently if directory exists (EXIST).
  - No explicit error for permission issues (logs success anyway).
- Edge Cases:
  - Invalid Path: Caller validates path starts with S1\_DIR, preventing issues like ...
  - Empty Path: Not expected, as path includes S1\_DIR.

- Security:
  - Ensures directories are within server scope by skipping S1\_DIR.
- Viva Tips:
  - Explain recursive creation for nested paths (e.g., ./S1/dir/subdir).
  - Note silent error handling as a design choice for simplicity.
- Example:
  - Input: create\_directory("./S1/dir/subdir").
  - Process: Creates ./S1/dir, then ./S1/dir/subdir.

## 5. list\_files(const char \*path, char \*file\_list, size\_t list\_size)

- Purpose: Lists files in a directory recursively, filtering by file type.
- Implementation:
  - Opens directory with opendir(path).
    - If fails (e.g., path invalid), sets file\_list to "ERROR: Directory not found" and returns.
  - Iterates entries using readdir.
  - Skips . and .. with strcmp(entry->d\_name, ".") == 0.
  - Builds full path with snprintf(full\_path, MAX\_PATH, "%s/%s", path, entry->d\_name).
  - For regular files (entry->d\_type == DT\_REG):
    - S2: Checks .pdf with strstr(entry->d\_name, ".pdf").
    - S3: Checks .txt.
    - S4: Checks .zip.
    - S1: No filter, includes all files.
    - Adds path (without S1\_DIR) to file\_list using snprintf(file\_list + offset, list\_size - offset, "%s\n", full\_path + strlen(S1\_DIR) + 1).
  - For directories (DT\_DIR), recursively calls list\_files.
  - Tracks offset to prevent buffer overflow.
  - If no files found, sets file\_list to "No files found\n".
  - Closes directory with closedir.
- Libraries Used:
  - <dirent.h>: opendir, readdir, closedir, struct dirent.
  - <string.h>: snprintf, strcmp, strstr.
- Error Handling:
  - Handles invalid directories.
  - Truncates list if list\_size exceeded.
- Edge Cases:
  - Empty Directory: Returns "No files found\n".
  - Buffer Overflow: Stops adding files, potentially truncating list.
  - Symbolic Links: Skipped, as DT\_REG checks for regular files.
- Viva Tips:
  - Highlight server-specific filtering (e.g., S2 for .pdf only).
  - Discuss recursion for nested directories.
  - Mention buffer size limitation (4096 bytes).
- Example:
  - Input: list\_files("./S2/dir", file\_list, BUFFER\_SIZE).
  - Output: file\_list contains dir/file1.pdf\n.

## 6. tar\_files(const char \*filetype, const char \*output\_tar)

- Purpose: Creates a tar archive of files with a given extension.
  - Implementation:
    - Builds command with snprintf:
      - Format: find %s -name "\*.%s" -type f | tar -cvf %s -T - 2>/dev/null.
      - Example (S2): find ./S2 -name "\*.pdf" -type f | tar -cvf pdf.tar -T - 2>/dev/null.
    - Executes command using system.
    - Checks if output\_tar exists with access(output\_tar, F\_OK).
      - If fails (no files or error), creates empty file with fopen(output\_tar, "w") to avoid send\_file errors.
    - Logs "Tar created successfully" or "Tar creation failed or no files found".
  - Libraries Used:
    - <string.h>: snprintf.
    - <stdlib.h>: system.
    - <unistd.h>: access.
    - <stdio.h>: fopen, fclose.
  - Error Handling:
    - Suppresses tar errors with 2>/dev/null.
    - Creates empty file as fallback.
  - Edge Cases:
    - No Files: Results in empty or missing tar, handled by creating empty file.
    - Large Files: tar handles any size, no code-level limit.
  - Security:
    - Hardcoded filetype (pdf, txt, zip, c) prevents injection.
    - Operates within S\_DIR, ensuring safety.
  - Performance:
    - system call is heavyweight; in-memory tar could be faster.
  - Viva Tips:
    - Explain why tar was used (standard tool, simplifies archiving).
    - Discuss empty file workaround for robustness.
    - Suggest alternatives like libtar for efficiency.
  - Example:
    - Input: tar\_files("pdf", "pdf.tar").
    - Output: Creates pdf.tar with .pdf files.
-

## Functions Specific to s1.c

7. `send_to_server(const char *ip, int port, const char *command, char *response, size_t resp_size, int is_file_transfer, int client_socket)`

- Purpose: Sends a command to a secondary server (S2, S3, S4) and relays file data if needed.
- Implementation:
  - Creates socket with `socket(AF_INET, SOCK_STREAM, 0)`.
    - If fails, sets response to "ERROR: Socket failed", logs, and returns.
  - Sets 5-second send/receive timeouts with `setsockopt(SO_SNDFTIMEO, SO_RCVTIMEO)`.
  - Sets server address (`struct sockaddr_in`) with `serv_addr.sin_family = AF_INET`, `serv_addr.sin_port = htons(port)`, and `inet_pton(AF_INET, ip, &serv_addr.sin_addr)`.
  - Connects using `connect`.
    - If fails, sets response to "ERROR: Connection failed", logs, and closes socket.
  - For `uploaddf`, adjusts path:
    - Parses command with `sscanf(command, "%*s %s %s", filename, dest_path)`.
    - Strips `S1_DIR` using `snprintf(new_cmd, BUFFER_SIZE, "uploaddf %s %s", filename, dest_path + strlen(S1_DIR) + 1)`.
    - Else, copies command to `new_cmd`.
  - Sends command with `send(sock, new_cmd, strlen(new_cmd), 0)`.
  - If `is_file_transfer` (for `uploaddf`):
    - Reads "READY" with `read(sock, ready_signal, 5)`.
      - If not "READY", sets response to "ERROR: Server not ready", logs, and returns.
    - Reads file size from `client_socket` with `read(client_socket, &file_size, sizeof(off_t))`.
      - If fails, sets response to "ERROR: Failed to receive file size", logs, and returns.
    - Sends file size to server with `send(sock, &file_size, sizeof(off_t), 0)`.
    - Relays data in chunks:
      - Reads from `client_socket` with `read`.
      - Sends to server with `send`.
      - Tracks received bytes.
      - If read or send fails, sets appropriate error in response and returns.
  - Reads server response with `read(sock, response, resp_size - 1)`.
    - If `valread_resp > 0`, null-terminates response.
    - Else, sets "ERROR: No response from secondary server".
  - Closes socket with `close(sock)`.
  - Logs command, READY signal, file transfer, and response.
- Libraries Used:
  - `<sys/socket.h>`: `socket`, `connect`, `send`, `read`, `setsockopt`.
  - `<arpa/inet.h>`: `inet_pton`, `htons`.
  - `<sys/time.h>`: `struct timeval`.
  - `<string.h>`: `snprintf`, `strncpy`, `strcmp`.
  - `<errno.h>`: `strerror`.
- Error Handling:
  - Handles socket creation, connection, READY signal, file transfer, and response failures.
- Edge Cases:
  - Invalid `client_socket`: Assumed valid by caller.

- Server Down: Timeout triggers error response.
- Empty File: Relays file\_size = 0 correctly.
- Security:
  - Path adjustment ensures correct directory mapping.
- Viva Tips:
  - Emphasize S1's proxy role, hiding S2–S4 from clients.
  - Discuss timeouts for robustness.
  - Explain path adjustment for uploaddf.
- Example:
  - Input: send\_to\_server("127.0.0.1", 9087, "uploaddf file.pdf dir", response, BUFFER\_SIZE, 1, client\_sock).
  - Output: Forwards file to S2, sets response to "SUCCESS: File uploaded".

#### **8. validate\_path(const char \*path)**

- Purpose: Ensures a path starts with S1\_DIR for security.
- Implementation:
  - Checks path != NULL and compares prefix with strncmp(path, S1\_DIR, strlen(S1\_DIR)) == 0.
  - Returns 1 if valid, 0 otherwise.
- Libraries Used:
  - <string.h>: strncmp, strlen.
- Error Handling:
  - Handles NULL path.
- Edge Cases:
  - NULL Path: Returns 0.
  - Short Path: If path is shorter than S1\_DIR, returns 0.
- Security:
  - Prevents access outside ./S1 (e.g., /etc).
- Viva Tips:
  - Highlight as a simple security measure.
  - Discuss importance of sandboxing in file servers.
- Example:
  - Input: validate\_path("./S1/dir/file.c") → 1.
  - Input: validate\_path("./S2/dir") → 0.

#### **9. prclient(int client\_socket)**

- Purpose: Processes client commands in a child process.
- Implementation:
  - Loops to read commands with read(client\_socket, buffer, BUFFER\_SIZE - 1).
    - If valread <= 0 (disconnect), logs "Client disconnected" and breaks.
  - Null-terminates buffer: buffer[valread] = '\0'.
  - Parses with sscanf(buffer, "%s %s %s", command, arg1, arg2), stores number of args.
  - Logs command.
  - Handles commands:
    - uploaddf (needs 3 args):
      - Validates arg2 with validate\_path.
      - For .c files (strstr(arg1, ".c")):
        - Calls create\_directory(arg2).
        - Calls receive\_file(client\_socket, filepath) with filepath = S1\_DIR/arg2/arg1.

- For others:
      - Sends "READY" to client.
      - For .pdf: Calls `send_to_server(S2_IP, PORT_S2, buffer, response, BUFFER_SIZE, 1, client_socket)`.
      - For .txt: Uses S3.
      - For .zip: Uses S4.
      - Else: Sends "ERROR: Invalid file type".
      - Sends response to client.
  - `downlf` (needs 2 args):
    - Validates arg1.
    - For .c files: Calls `send_file(client_socket, arg1)`.
    - For others:
      - Adjusts path (e.g., `./S1/dir/file.pdf` → `./S2/dir/file.pdf`).
      - Connects to server (S2, S3, or S4).
      - Sends `downlf adjusted_path`.
      - Checks "READY" and `file_size`.
      - Relays file to client with `send(client_socket, "READY", 5, 0)`, sends size, and forwards data chunks.
      - Sends error if READY/size fails.
  - `removef` (needs 2 args):
    - Validates arg1.
    - For .c files: Calls `remove(arg1)`, sends "SUCCESS: File removed" or "ERROR: Removal failed".
    - For others: Adjusts path, calls `send_to_server` with `removef`, sends response.
  - `downltar` (needs 2 args):
    - For c: Calls `tar_files("c", "cfiles.tar")`, sends file, removes tar.
    - For pdf, txt, zip:
      - Calls `send_to_server` to S2, S3, or S4.
      - If response is "READY", connects again, sends command, calls `receive_file(client_socket, "pdf.tar")` (or `text.tar, zip.tar`).
      - Else, sends response (e.g., "ERROR: No files to tar").
      - Else: Sends "ERROR: Invalid filetype".
  - `dispfnames` (needs 2 args):
    - Validates arg1.
    - Initializes `file_list` with "Files in directory:\n".
    - Calls `list_files(arg1)` for local files.
    - Queries S2, S3, S4 with adjusted paths (e.g., `./S2/dir`).
    - Combines responses, ignoring errors.
    - If no files, sets "No files found\n".
    - Sends `file_list` to client.
  - Else: Sends "ERROR: Invalid command".
- Libraries Used:
    - `<sys/socket.h>`: `read`, `send`, `socket`, `connect`.
    - `<string.h>`: `sscanf`, `strcmp`, `snprintf`, `strstr`.
    - `<arpa/inet.h>`: `inet_pton`, `htons`.
    - `<unistd.h>`: `usleep`, `close`.
    - `<stdlib.h>`: `system` (via `tar_files`).

- Error Handling:
  - Invalid paths, file types, commands, and disconnects.
  - Server connection failures for forwarded commands.
- Edge Cases:
  - Partial Tar Download: Handled by TCP reliability.
  - Large File Lists: Limited by BUFFER\_SIZE \* 4, may truncate.
  - Invalid Args: Checks args count.
- Concurrency:
  - Runs in child process, allowing multiple clients.
- Viva Tips:
  - Explain S1's delegation model.
  - Discuss fork-based concurrency vs. threading.
  - Highlight unified client interface.
- Example:
  - Input: uploadf file.c ./S1/dir.
  - Output: Saves file, sends "SUCCESS: File uploaded".

#### **10. main() in s1.c**

- Purpose: Sets up and runs S1 server.
- Implementation:
  - Creates S1\_DIR with mkdir(S1\_DIR, 0777).
  - Initializes log file with fopen(LOG\_FILE, "a"), closes it.
  - Creates socket with socket(AF\_INET, SOCK\_STREAM, 0).
    - If fails, logs and exits.
  - Sets SO\_REUSEADDR with setsockopt to allow port reuse.
  - Binds to port 9086 with bind.
    - If fails, logs and exits.
  - Listens with listen(server\_fd, 10).
    - If fails, logs and exits.
  - Logs and prints "S1 started, listening on port 9086".
  - Loops:
    - Accepts clients with accept.
      - If fails, logs and continues.
    - Logs client IP/port using inet\_ntop.
    - Forks child process:
      - Child: Closes server\_fd, calls prcclient(client\_socket), exits.
      - Parent: Closes client\_socket.
    - Cleans zombies with waitpid(-1, NULL, WNOHANG).
  - Closes server\_fd (unreachable).
- Libraries Used:
  - <sys/socket.h>: socket, bind, listen, accept, setsockopt.
  - <arpa/inet.h>: inet\_ntop, htons.
  - <sys/stat.h>: mkdir.
  - <stdio.h>: printf, fopen, fclose.
  - <sys/wait.h>: waitpid.
  - <unistd.h>: fork, close.
- Error Handling:
  - Socket setup and accept failures.
- Edge Cases:

- Fork Failure: Not handled (rare, process limit).
    - Port in Use: SO\_REUSEADDR mitigates.
  - Viva Tips:
    - Discuss fork model for concurrency.
    - Explain SO\_REUSEADDR for quick restarts.
    - Highlight zombie cleanup.
  - Example:
    - Output: "S1 server listening on port 9086".
    - Process: Accepts clients, forks handlers.
- 

## Functions in s2.c, s3.c, s4.c

### 11. main() in s2.c, s3.c, s4.c

- Purpose: Sets up and runs S2, S3, or S4.
- Implementation:
  - Creates directory (S2\_DIR, S3\_DIR, S4\_DIR) with mkdir.
  - Initializes log file.
  - Creates socket with socket(AF\_INET, SOCK\_STREAM, 0).
    - If fails, logs and exits.
  - Binds to port (9087, 9088, 9089).
    - If fails, logs and exits.
  - Listens with listen(server\_fd, 10).
    - If fails, logs and exits.
  - Logs and prints "S2 listening on port 9087" (or S3/S4).
  - Loops:
    - Accepts connections with accept.
      - If fails, logs and continues.
    - Logs "New connection from S1".
    - Reads command with read(client\_socket, buffer, BUFFER\_SIZE - 1).
      - If valread <= 0, closes socket and continues.
    - Parses with sscanf(buffer, "%s %s %s", command, arg1, arg2).
    - Handles commands:
      - uploadf: Calls create\_directory(dirpath) (S2\_DIR/arg2), receive\_file(client\_socket, filepath) (S2\_DIR/arg2/arg1).
      - downlf: Validates arg1 starts with S2\_DIR, calls send\_file.
      - removef: Validates arg1, calls remove, sends "SUCCESS" or "ERROR".
      - downltar: Calls tar\_files("pdf", "pdf.tar") (or txt, zip), sends file, removes tar.
      - dispfnames: Calls list\_files(arg1, file\_list, BUFFER\_SIZE), sends list.
    - Closes client\_socket.
  - Closes server\_fd (unreachable).
- Libraries Used:
  - <sys/socket.h>: socket, bind, listen, accept.
  - <arpa/inet.h>: htons.
  - <sys/stat.h>: mkdir.
  - <stdio.h>: printf, fopen, fclose.
  - <string.h>: sscanf, strcmp, snprintf.
- Error Handling:

- Socket setup and command read failures.
  - Edge Cases:
    - Buffer Overflow: Truncates at BUFFER\_SIZE - 1, unlikely for commands.
    - Single Client: Handles one S1 request at a time.
  - Viva Tips:
    - Note specialization (S2 for .pdf).
    - Discuss single-threaded design (sufficient for S1).
  - Example:
    - Output: "S2 listening on port 9087".
    - Process: Handles .pdf requests.
- 

## Functions in w25clients.c

12. send\_file(int socket, const char \*filepath)

- Purpose: Sends a file to the server.
- Implementation:
  - Opens file with open(filepath, O\_RDONLY).
    - If fails, prints error with strerror, sends file\_size = -1, returns.
  - Gets size with fstat(fd, &st).
  - Sends size with send(socket, &file\_size, sizeof(off\_t), 0).
    - If fails, prints error and returns.
  - Reads chunks with read(fd, buffer, BUFFER\_SIZE) and sends with send.
    - If read or send fails, prints error.
  - Closes file with close(fd).
  - Prints "File sent successfully".
- Libraries Used:
  - <fcntl.h>: open, read, close.
  - <sys/stat.h>: fstat.
  - <sys/socket.h>: send.
  - <stdio.h>: printf.
  - <errno.h>: strerror.
- Error Handling:
  - File open, read, and send failures.
- Edge Cases:
  - Empty File: Sends file\_size = 0.
- Viva Tips:
  - Note protocol consistency with servers.
  - Discuss error feedback for usability.
- Example:
  - Input: send\_file(sock, "file.c").
  - Output: "File sent successfully".

13. receive\_file(int socket, const char \*filename)

- Purpose: Receives a file and saves it locally.
- Implementation:
  - Reads "READY" with read(socket, ready\_signal, 5).
    - If not "READY", reads error message, prints, and returns.
  - Reads size with read(socket, &file\_size, sizeof(off\_t)).
    - If fails, prints error and returns.

- If `file_size == -1`, reads and prints error message, returns.
  - Opens file with `open(filename, O_WRONLY | O_CREAT | O_TRUNC, 0666)`.
    - If fails, prints error and returns.
  - Reads chunks with `read` and writes with `write`.
    - Tracks received bytes.
    - If `write` fails, prints error and returns.
  - Closes file with `close(fd)`.
  - Prints "SUCCESS: File downloaded as filename" if received == `file_size`, else "ERROR: Incomplete file download".
- Libraries Used:
  - `<sys/socket.h>`: `read`.
  - `<fcntl.h>`: `open`, `write`, `close`.
  - `<string.h>`: `strcmp`.
  - `<stdio.h>`: `printf`.
  - `<errno.h>`: `strerror`.
- Error Handling:
  - `READY`, `size`, and `write` failures.
- Edge Cases:
  - Empty File: Handled correctly (`file_size = 0`).
- Viva Tips:
  - Highlight transfer verification.
  - Discuss error messages for user clarity.
- Example:
  - Input: `receive_file(sock, "file.c")`.
  - Output: Saves file, prints "SUCCESS: File downloaded as file.c".

#### 14. `read_full_response(int socket, char *buffer, size_t size)`

- Purpose: Reads a response until newline or buffer limit.
- Implementation:
  - Reads with `read(socket, buffer + total_read, size - total_read - 1)`.
  - Loops until newline (`strchr(buffer, '\n')`), buffer full, or no data.
  - Null-terminates buffer.
  - Returns bytes read.
- Libraries Used:
  - `<sys/socket.h>`: `read`.
  - `<string.h>`: `strchr`.
- Error Handling:
  - Handles no data (timeout).
- Edge Cases:
  - No Newline: Reads until `size - 1`.
  - Empty Response: Returns 0.
- Viva Tips:
  - Note use for text responses (`dispfnames`, `errors`).
  - Contrast with `receive_file` for binary data.
- Example:
  - Input: `read_full_response(sock, buffer, BUFFER_SIZE)`.
  - Output: `buffer` has "SUCCESS: File removed\n".

#### 15. `main()` in `w25clients.c`

- Purpose: Runs client interface.

- Implementation:
  - Prints "Welcome to w25clients".
  - Loops:
    - Prompts "w25clients\$".
    - Reads command with fgets(command, BUFFER\_SIZE, stdin).
      - Removes newline with strcspn(command, "\n").
      - Breaks on EOF.
    - Creates socket with socket(AF\_INET, SOCK\_STREAM, 0).
    - Sets 5-second timeouts with setsockopt.
    - Connects to S1 (127.0.0.1:9086).
    - Parses command with sscanf(command, "%s %s %s", cmd, arg1, arg2).
    - Handles commands:
      - uploadf (3 args):
        - Checks file exists with stat(arg1, &st).
        - Sends command.
        - Waits for "READY".
        - Calls send\_file(arg1).
        - Reads and prints response.
      - downlf (2 args):
        - Sends command.
        - Extracts filename with strrchr(arg1, '/').
        - Calls receive\_file(socket, filename).
      - removef (2 args):
        - Sends command.
        - Reads and prints response.
      - downltar (2 args):
        - Validates filetype (c, pdf, txt, zip).
        - Sends command.
        - Sets filename (cfiles.tar, pdf.tar, etc.).
        - Calls receive\_file.
        - Verifies file with stat, prints size.
      - dispfnames (2 args):
        - Sends command.
        - Reads and prints response.
      - exit: Closes socket, prints "Exiting client", breaks.
      - Else: Prints help message with valid commands.
    - Closes socket.
- Libraries Used:
  - <stdio.h>: printf, fgets.
  - <string.h>: sscanf, strcmp, strrchr, strcspn.
  - <sys/socket.h>: socket, connect, send, read, setsockopt.
  - <arpa/inet.h>: inet\_pton, htons.
  - <sys/stat.h>: stat.
  - <sys/time.h>: struct timeval.
  - <errno.h>: strerror.
- Error Handling:
  - Socket, connection, file, and command errors.
- Edge Cases:

- Invalid Filetype: Prints error for download.
  - EOF Input: Exits gracefully.
  - Viva Tips:
    - Highlight user-friendly interface.
    - Discuss abstraction of distributed system.
  - Example:
    - Input: uploadf file.c ./S1/dir.
    - Output: Sends file, prints "SUCCESS: File uploaded".
- 

## Viva Preparation and Additional Insights

### Key Points to Emphasize

1. Distributed Design:
  - S1 coordinates, S2–S4 specialize, mimicking cloud storage (e.g., AWS S3).
  - S1's proxy role simplifies client interactions.
2. Socket Programming:
  - Uses TCP (SOCK\_STREAM, AF\_INET) for reliability.
  - S1 forks for concurrency; S2–S4 are single-threaded for simplicity.
3. File Operations:
  - READY-size-data protocol ensures robust transfers.
  - Hierarchical directories (./S1/dir) support organization.
  - Tar archiving uses system for simplicity.
4. Security:
  - Path validation prevents unauthorized access.
  - Hardcoded filetypes avoid command injection.
5. Robustness:
  - 5-second timeouts prevent hangs.
  - Detailed error messages aid debugging.
  - Logs capture all actions for tracing.
6. Scalability:
  - S1's forking supports multiple clients but has memory overhead.
  - S2–S4 could be scaled with threading if direct client access is added.

### Potential Viva Questions

- Q: How does S1 route files?
  - A: Checks file extension in prclient: .c stays local, .pdf to S2, .txt to S3, .zip to S4, using send\_to\_server.
- Q: What if a server is down?
  - A: send\_to\_server times out after 5 seconds, sends "ERROR: Connection failed" to client, logged for debugging.
- Q: Why use system for tar?
  - A: Simplifies archiving vs. parsing tar format manually. Trade-off is performance; libtar could improve efficiency.
- Q: How are large files handled?
  - A: Chunked transfers (4096 bytes) in send\_file/receive\_file handle any size, limited only by disk.
- Q: Security measures?
  - A: validate\_path ensures files stay in server directories. Commands are parsed to prevent misuse.

### Improvements to Suggest

- Use threading instead of forking in S1 for lower memory use.
- Buffer logs or keep file open for efficiency.
- Dynamic buffer sizes for list\_files to avoid truncation.
- Add authentication for client-server interactions.